

PBA Software Development
Parallel Programming

Concurrency

Concepts, problems, and solutions



Units of execution

- Reasoning about concurrency and parallelism requires us to have terms to discuss that which is executing.
- Two common types of execution unit:
 - Processes
 - Threads

Processes

- Common OS abstraction. Represents a running program.
- Private address space relative to other processes.
- Contextual information such as register set, IO handles, etc...

Threads

- A process contains one or more threads.
- Threads have private register set and stack.
- Threads within a process share memory with each other.
 - Threads may also have private per-thread “thread local memory”.

Modern operating systems


- The line between threads and processes can be blurry.
 - Often both are implemented on top of the same OS abstraction, *kernel threads*.
- The critical distinction for many discussions is data sharing.
 - Processes do not share address space without explicit assistance.
 - Threads within a process share address space.

Terminology

- In the concurrent programming world, different projects use different terminology for execution units.
 - Threads
 - Tasks
 - Processes
 - Activities
- Generally a good idea to look up precisely what a term means in whatever context you encounter it in.

Parallel vs. concurrent

- What is the difference between a parallel and a concurrent program?



Parallel vs. concurrent


- Consider two units of execution that are started at the same time.
- If they run in **parallel**, then over any time interval during their execution, both may be executing at the same time.
- If they run **concurrently**, they *may* execute in parallel, or they may be sequentialized where only one makes progress at any given point in time and their execution is interleaved.
 - Multitasking operating systems have exploited this for many years.

Parallel vs. concurrent

- Consider three programs executing concurrently.



Concurrent, non-parallel execution



Concurrent, parallel execution

Shared vs. distributed memory

- Common distinction when considering parallel and networked computers.
- Consider a memory address space and a set of processing elements (CPUs or single CPU cores).

Shared memory

- Every processing unit can directly address every element of the address space.
- Examples include:
 - Multicore CPUs
 - Each core can see all of the RAM in the machine.
 - Shared memory parallel systems
 - Traditional parallel computers, such as servers.

Distributed memory

- Processing elements can directly address only a proper subset of the address space.
 - In order to access memory outside of this subset, processing elements must either:
 - Be assisted by another processing element that can directly address the memory.
 - Use specialized mechanisms (such as network hardware) to access the memory.
- We distinguish directly from indirectly accessible memory as **local** versus **remote**.

Dependencies

- Dependencies are a critical property of any system that impacts concurrency.
 - Even impacts real-world activities like cooking.
- A dependency is a state (either of data or control) that must be reached before a part of a program can execute.

Simple example

- Both x and y must be defined before they can be read and used in the computation of z.
- Therefore z **depends on** the assignment statements that associate values with x and y.

```
x=6;  
y=7;  
z=x+y;
```

Dependencies limit parallelism

- If part of a program (a “subprogram”) depends on another subprogram, then the dependent portion must wait until that which it depends on is complete.
 - The dependency dictates sequentialization of the subprograms.
- If no such dependency exists, the subprograms can execute in parallel (“embarrassingly” parallel).

Bernstein's conditions

- We can formalize this via Bernstein's conditions.
- Consider a subprogram P .
- Let $IN(P)$ represent the set of memory locations (including registers) or variables that P uses as input by reading from them.
- Let $OUT(P)$ represent a similar set of locations that P uses as output by writing to them.
- We can use these sets to determine if two subprograms P_1 and P_2 are dependent, and therefore whether or not they can execute in parallel.

Bernstein's conditions

- Given two subprograms P_1 and P_2 , their execution in parallel is equivalent to their execution in sequence if the following conditions hold.

$$BC1 \quad : \quad OUT(P_1) \cap OUT(P_2) = \emptyset$$

$$BC2 \quad : \quad IN(P_1) \cap OUT(P_2) = \emptyset$$

$$BC3 \quad : \quad IN(P_2) \cap OUT(P_1) = \emptyset$$

Atomicity

- Atom derived from Greek word for *indivisible*.
- Often we write programs that include sequences of operations that we would like to behave as though they were a single indivisible operation.
- Classic example:
Two units of execution updating a shared variable

Intermediate results

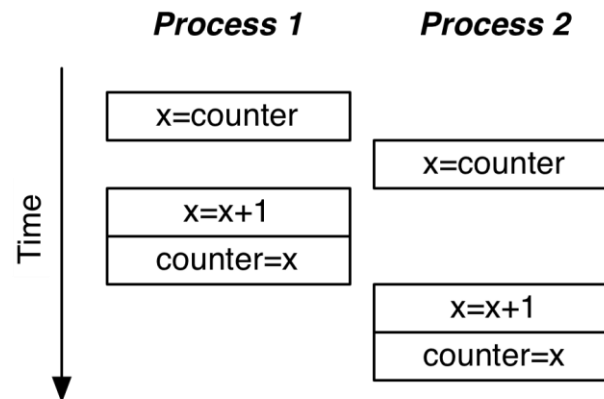
- Consider a shared counter.

```
shared int counter;  
private int x;  
x = counter;  
x = x+1;  
counter = x;
```

- We assume incrementing the counter is an atomic operation. If we do not, we could miss updates.

Missed update

- Poor interleaving of concurrent threads of execution can result in missed updates if atomicity isn't enforced.



Critical section

- Sequences of operations that can reveal invalid or intermediate data are referred to as *critical sections*.
- In the counter example, not only would this mean the increment operator would be a critical section, but...
 - Any other operation that modifies the counter would be.
 - As would any operation reading the counter.

Mutual exclusion

- Mutual exclusion refers to ensuring that one and only one thread of execution be in a critical section at any point in time.
- Concurrency control mechanisms can be used to implement this.

Thread safety

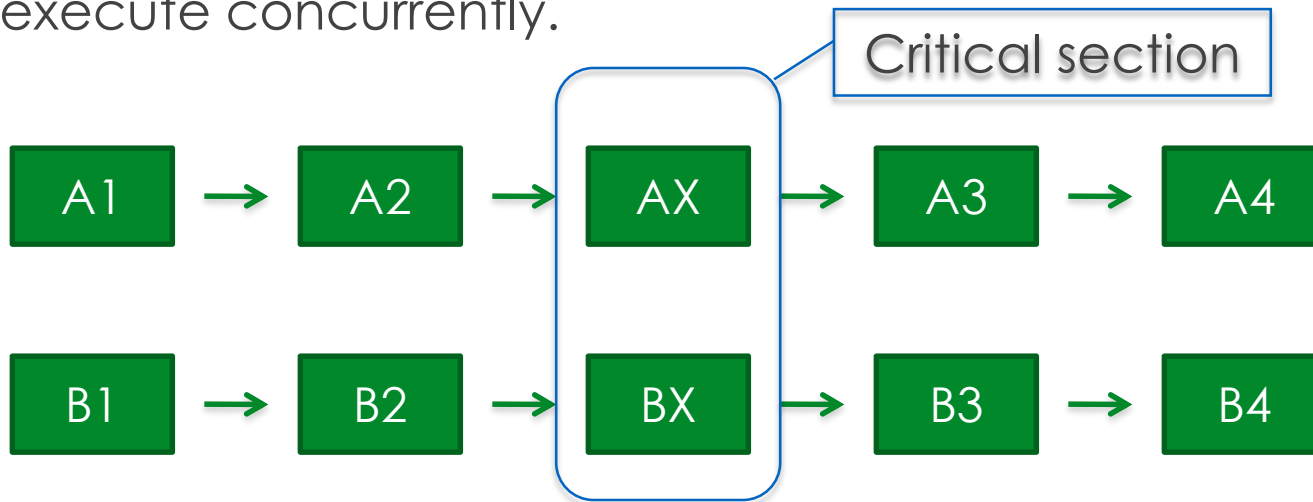
- A common term to encounter in programming is *thread safety*.
- Typically used to describe how a programmer can expect a subroutine, object, or data structure to behave in a concurrent context.
- A thread safe implementation will behave as expected in the presence of multiple threads of execution, that is, the same behavior as if done sequentially.

Reentrant code

- A related term is *reentrant*.
- A reentrant routine (or region of code) may be entered by a thread of execution even if another thread has already started executing it at the same time.
- Non-reentrant code typically relies on state, such as global variables, that are not intended to be shared by multiple threads of execution.

Race conditions

- Consider two sequences of operations that can execute concurrently.



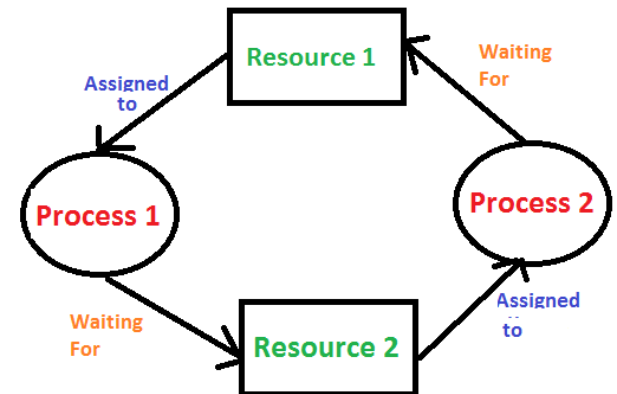
- A race condition exists if the result of the sequences depends on their relative arrival at some critical point in the sequence.
 - Indicated here as AX and BX.

Race conditions

- If the result of operations after the critical point in the sequence (AX or BX) is dependent on which executes first (AX then BX versus BX then AX), a race is present in the code.
 - The result is dependent on which thread of execution “wins” the race to that point.
- Code without a race means that either:
 - The relative ordering doesn't matter.
 - Concurrency control constructs are used to enforce the required ordering.

Deadlock

- Deadlock results from a cyclical chain of dependencies.
 - A depends on B, and B depends on A.
- Often results from a misuse of concurrency control constructs.
 - A holds lock on resource that B wants, B holds lock on resource that A wants.
- Possible to observe when a program “freezes”, and simply ceases to make progress.



Livelock

- Similar to deadlock, but instead of freezing, a program enters into an endless cycle of operations that it cannot continue past.
 - Example: Acquire lock 1, attempt to acquire lock 2 and fail, release lock 1, retry.
- Livelock is detectable, but can be less obvious from simple observation than deadlock due to the fact that the program counter doesn't freeze.
 - It does enter a repeating cycle though.

Liveness

- If multiple threads of execution are started, we expect that they will all eventually finish.
- Often we will go further, and assume that during any reasonable interval of time (such as 1 second), all concurrent threads will make some progress.
- When a system or program ensures that this property holds, we say it ensures liveness.

Liveness

- Often liveness is a property of a scheduler, either within the OS, thread library, or user application.
- Concurrent programmers that manage their own threads must occasionally consider liveness of their management algorithms.
- When a system doesn't ensure liveness, then starvation results.

Starvation

- A thread of execution will starve if it ceases to make progress in the presence of others.
 - Example: a program that is continuously preempted by others and never allowed to execute.
- Priority and aging schemes can be used to ensure that starvation does not occur due to a scheduler.
- Queuing disciplines (such as FIFO) can guarantee starvation will not occur due to synchronization primitives.

Nondeterminism

- Not always a problem – some algorithms can exploit nondeterminism to their advantage.
 - *Example:* concurrent queries to multiple search engines, return whichever result returns first.
- Nondeterminism is an inherent property of parallel systems, and must be kept in mind at all times.
 - Both for positive and negative reasons.

Control

- To prevent correctness problems, we must control how concurrently executing programs interact.
- Synchronization
 - Mutual exclusion
 - Semaphores, locks, monitors
- Transactions
 - Speculative execution and conflict resolution

Synchronization

- Synchronization allows concurrent threads of execution to communicate and coordinate information about their relative state.
- This can be used to coordinate their execution and prevent correctness problems.

Semaphores

- Semaphores are one of the original synchronization primitives due to Dijkstra.
- Represented as:
 - A single integer.
 - A queue of blocked accessors.
- Two operations are used to manipulate semaphores.

Semaphores

- Consider a semaphore **s**.
 - Initialized to 1 with an empty queue.
- **P(s)**:
 - Executed just before entering a critical section
 - If $s > 0$ then decrement s ,
Otherwise, block the thread that attempted to perform P(s).
- **V(s)**:
 - Executed just before leaving a critical section
If another thread is blocked on a call to P(s), unblock it.
Otherwise, increment s .
- Increment/decrement apply **atomically** to the integer.
- Collection used to manage set of blocked threads.

Semaphores

- Binary Semaphore
 - Initial value is 1.
 - Allows one process in the critical section at a time (Mutual Exclusion).
- Counting Semaphore
 - Initial value is N.
 - The semaphore allows up to N processes in the critical section at the same time.
- Semaphores are a general structure.

Semaphores in .Net

- `System.Threading.Semaphore`
 - Counting semaphore
 - With no name: Local semaphore
 - With name: System semaphore (cross process).
- Methods:
 - `WaitOne`: Blocking wait to enter Critical Section
 - `Release`: Release the Critical Section.

Semaphores in .Net

- `System.Threading.SemaphoreSlim`
 - Counting semaphore
 - Local semaphore
- Methods:
 - **Wait**: Blocking wait to enter Critical Section
 - **WaitAsync**: Non-blocking wait for Critical Section. Returns a task containing Critical Section code.
 - **Release**: Release the Critical Section.

Locks

- A lock is a familiar synchronization construct ensuring mutual exclusion.
 - Consider locks on doors that have only one key.
 - Someone obtains key, is able to unlock (“acquire”) the lock.
 - Others must wait for the key to be released before they can unlock the lock themselves.
- Locks can be implemented using binary semaphores.

Locks in .Net

- `lock(lockObject) { ... }`
 - Local mutual exclusion.
 - Any object may be used as "key".
 - Performance overhead to lock and unlock: 20 ns.
- **Mutex**
 - Cross-process mutual exclusion.
 - Performance overheads to lock and unlock: 1000 ns.
 - Methods:
 - **WaitOne**: Blocking wait to enter critical section.
 - **ReleaseMutex**: Releasing critical section.

Locks in .Net

■ ReaderWriterLockSlim

- Local Locking protocol.
- Allows more readers simultaneously (shared lock) but only one writer at a time (exclusive lock)
- Performance overheads to lock and unlock: 40 ns.
- Methods:
 - **EnterReadLock**
 - **ExitReadLock**
 - **EnterWriteLock**
 - **ExitWriteLock**

Monitors

- Object-oriented approach to synchronization.
- Encapsulates the synchronization structures along with the resource to be guarded, and makes the resource **Thread-Safe**.
- Better than leaving the synchronization problem to the developers.

Monitors

- A binary semaphore is used to enforce mutual exclusion of the monitor, that is, only one process at a time will be using the monitor.

Monitors – Condition variable

- Important feature of monitors beyond encapsulation is introduction of *condition variables*.
- Condition variables allow threads to signal each other.
 - Beyond simple lock semantics.

Monitors

- Condition variable signaling
 - A thread may be executing within a monitor routine and reach a state where it must block and yield the monitor to allow another thread to access it so it can later pick up where it left off and continue.
 - Two operations are provided for condition variables.
 - Wait()
 - Signal()

Monitors

- A thread may *wait* on a condition variable.
 - Yield the synchronization structures for the monitor to allow other threads in.
 - Block waiting for another thread to signal it.
- Another thread may later *signal* threads blocked and waiting on a condition variable.
 - Signaled threads will wake up and attempt to reacquire the synchronization structures of the monitor.
 - When they succeed, they pick up immediately after the point where they *wait()*-ed.

Monitor in .Net

- `System.Threading.Monitor`
 - Used to ensure mutual access to a data resource.
 - Performance overhead: 200 ns
 - Methods:
 - `Enter`: Blocking wait to enter the monitor.
 - `Exit`: releasing the monitor.
 - `Wait(object)`: blocking wait for a condition object.
 - `Pulse(object)`: Signalling release on a condition object.

Further reading

- http://www.albahari.com/threading/part2.aspx#_Locking